

Chiem MAT 115 Homework 3

For the third homework assignment, we will use the **babynames** dataset. This is a dataset of US baby names using information provided from the US Social Security Administration (SSA). As per usual, you should write your answers in RMarkdown. The due date for this assignment is **Sept. 25 at 2:30 PM**.

1. The **babynames** dataset is in its own package. Download it and load it into R. Display the first 8 rows of the dataset.

Note: the **babynames** dataset is a **tibble**, which is a different format from the traditional data frame. For our purposes here, it does not influence our use of the data.

```
#install.packages('babynames')
library(babynames)
babynames
```

```
## # A tibble: 1,924,665 x 5
##   year sex  name      n  prop
##   <dbl> <chr> <chr>   <int> <dbl>
## 1  1880 F    Mary    7065 0.0724
## 2  1880 F    Anna    2604 0.0267
## 3  1880 F    Emma    2003 0.0205
## 4  1880 F  Elizabeth 1939 0.0199
## 5  1880 F   Minnie   1746 0.0179
## 6  1880 F  Margaret 1578 0.0162
## 7  1880 F    Ida     1472 0.0151
## 8  1880 F   Alice   1414 0.0145
## 9  1880 F  Bertha   1320 0.0135
##10  1880 F   Sarah   1288 0.0132
## # i 1,924,655 more rows
```

2. Use your R indexing skills to find out how popular (in terms of **prop**) your given name was in the year you were born. If your name is not in the dataset, then just pick a name that interests you for whatever reason.

```
babynames[babynames$name == "Damien" & babynames$year == 2004, ]
```

```
## # A tibble: 2 x 5
##   year sex  name      n  prop
##   <dbl> <chr> <chr>   <int> <dbl>
## 1  2004 F    Damien    11 0.00000545
## 2  2004 M    Damien   1955 0.000926
```

```
#length(babynames[babynames$name == "Damien" & babynames$year == 2004, ])
```

A: for males my name had a prop of 0.00092558 and 0.00000545 for females in my birth year.

3. Subset the dataset into two separate ones by `sex`. For each sex, take the top-25 names by `n`. Do not over think this request. Duplicates of a name across years is acceptable. Make a boxplot of `n` by `name`. Beautify these plots in at least two ways. What strikes you about these plots? How are they similar and how are they different?

```
maleNames <- babynames[babynames$sex == "M",]

femaleNames <- babynames[babynames$sex == "F",]

maleNames <- maleNames[order(maleNames$n, decreasing = TRUE)[1:25],]
femaleNames <- femaleNames[order(femaleNames$n, decreasing = TRUE)[1:25],]

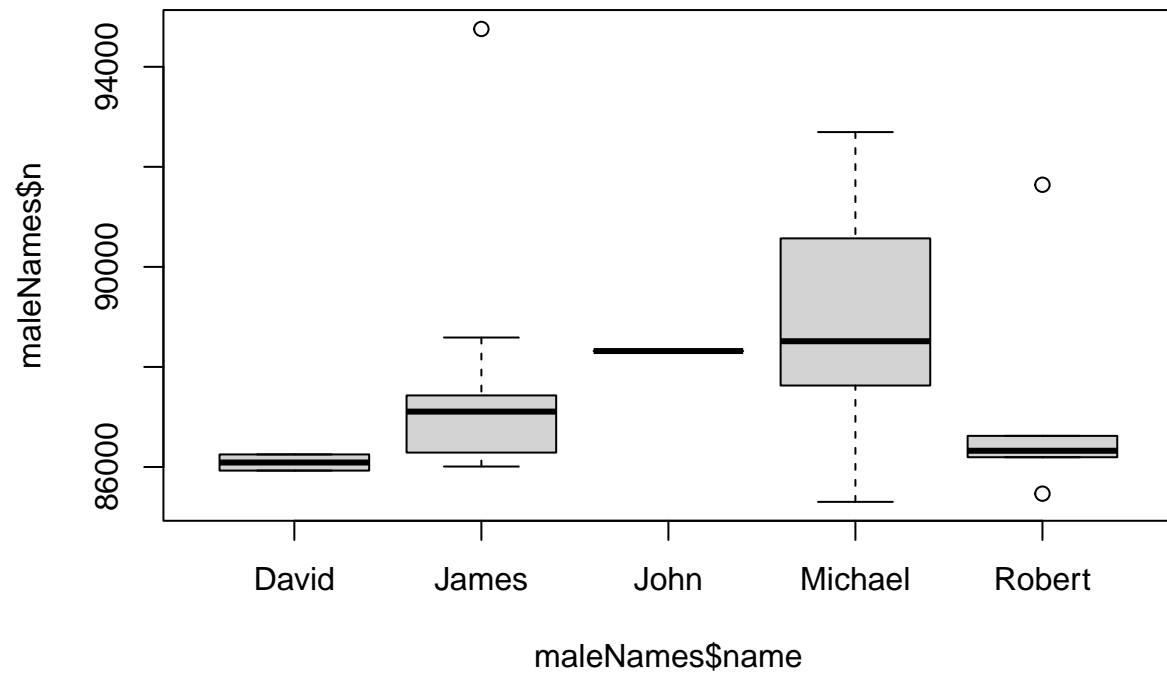
maleNames
```

```
## # A tibble: 25 x 5
##   year sex  name      n  prop
##   <dbl> <chr> <chr>   <int> <dbl>
## 1  1947 M    James  94756 0.0510
## 2  1957 M   Michael 92695 0.0424
## 3  1947 M   Robert 91642 0.0493
## 4  1956 M   Michael 90620 0.0423
## 5  1958 M   Michael 90520 0.0420
## 6  1948 M    James  88588 0.0497
## 7  1954 M   Michael 88514 0.0428
## 8  1955 M   Michael 88335 0.0423
## 9  1947 M    John   88318 0.0475
## 10 1946 M    James  87431 0.0530
## # i 15 more rows
```

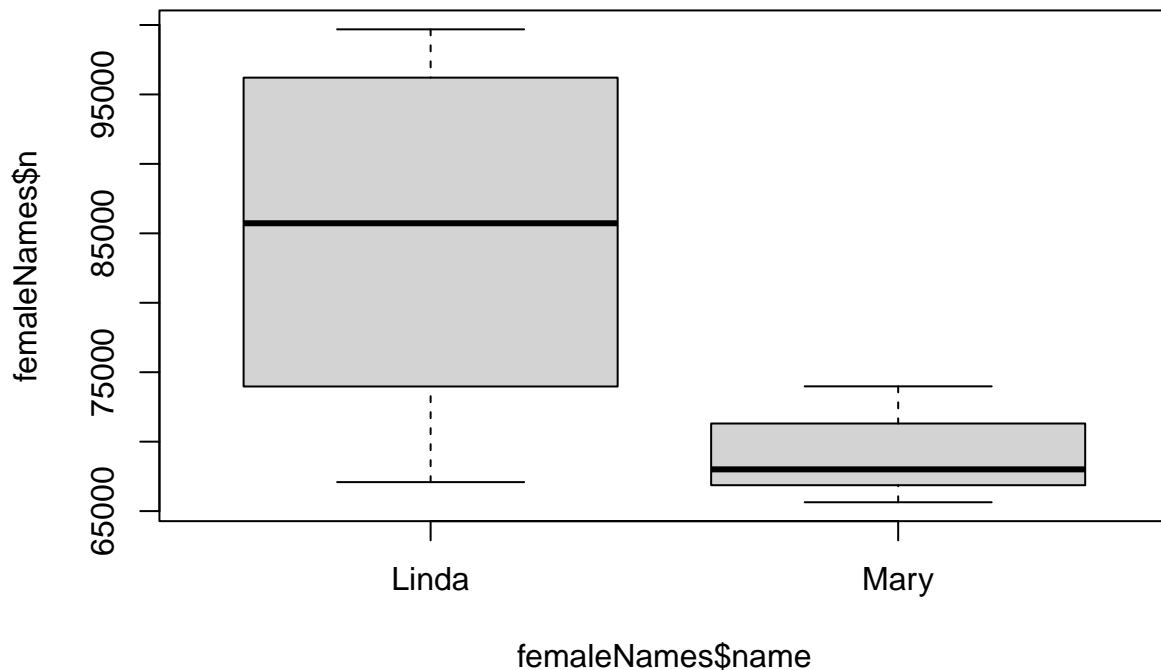
```
femaleNames
```

```
## # A tibble: 25 x 5
##   year sex  name      n  prop
##   <dbl> <chr> <chr>   <int> <dbl>
## 1  1947 F    Linda 99686 0.0548
## 2  1948 F    Linda 96209 0.0552
## 3  1949 F    Linda 91016 0.0518
## 4  1950 F    Linda 80432 0.0457
## 5  1921 F    Mary  73982 0.0578
## 6  1951 F    Linda 73972 0.0400
## 7  1924 F    Mary  73532 0.0568
## 8  1922 F    Mary  72175 0.0579
## 9  1947 F    Mary  71688 0.0394
## 10 1923 F    Mary  71635 0.0572
## # i 15 more rows
```

```
boxplot(maleNames$n ~ maleNames$name)
```



```
boxplot(femaleNames$n ~ femaleNames$name)
```



4. For the entire `babynames` dataset, write a `for` loop that calculates the mean `prop` for each `year`. In your own words, explain what the output, sequence, and body of your code are doing.

```
propM <- vector(length=nrow(babynames[duplicated(babynames$year) == FALSE,]))
each_year <- vector(length=nrow(babynames[duplicated(babynames$year) == FALSE,]))
years <- babynames[duplicated(babynames$year) == FALSE, ]

for(i in 1:length(propM)){
  tempMatrix <- babynames[babynames$year == years$year[i],]

  propM[i] <- mean(tempMatrix$prop)
  each_year[i] <- tempMatrix$year[1]
}

#system.time(for(i in 1:length(propM)){
#  tempMatrix <- babynames[babynames$year == years$year[i],]

#  propM[i] <- mean(tempMatrix$prop)
#  each_year[i] <- tempMatrix$year[1]
#}) # 2.14 seconds

#propM
#each_year

propM_and_year <- data.frame(year = each_year, mean_prop = propM)
propM_and_year
```

##	year	mean_prop
## 1	1880	9.327296e-04
## 2	1881	9.615283e-04
## 3	1882	8.762491e-04
## 4	1883	8.952682e-04
## 5	1884	8.141704e-04
## 6	1885	8.139021e-04
## 7	1886	7.820985e-04
## 8	1887	7.867618e-04
## 9	1888	7.065885e-04
## 10	1889	7.226397e-04
## 11	1890	6.944554e-04
## 12	1891	7.028995e-04
## 13	1892	6.410328e-04
## 14	1893	6.616759e-04
## 15	1894	6.363961e-04
## 16	1895	6.141118e-04
## 17	1896	6.051479e-04
## 18	1897	6.169039e-04
## 19	1898	5.737293e-04
## 20	1899	6.126111e-04
## 21	1900	5.017828e-04
## 22	1901	5.907095e-04
## 23	1902	5.550258e-04
## 24	1903	5.499067e-04
## 25	1904	5.243434e-04
## 26	1905	5.102156e-04
## 27	1906	5.137318e-04
## 28	1907	4.736429e-04
## 29	1908	4.655718e-04
## 30	1909	4.425246e-04
## 31	1910	4.053799e-04
## 32	1911	3.868689e-04
## 33	1912	2.997358e-04
## 34	1913	2.740336e-04
## 35	1914	2.403288e-04
## 36	1915	2.056436e-04
## 37	1916	1.986359e-04
## 38	1917	1.943714e-04
## 39	1918	1.855249e-04
## 40	1919	1.858838e-04
## 41	1920	1.794639e-04
## 42	1921	1.779018e-04
## 43	1922	1.794109e-04
## 44	1923	1.814584e-04
## 45	1924	1.778298e-04
## 46	1925	1.817062e-04
## 47	1926	1.848431e-04
## 48	1927	1.858832e-04
## 49	1928	1.904784e-04
## 50	1929	1.971559e-04
## 51	1930	1.978530e-04

52 1931 2.082921e-04
53 1932 2.063539e-04
54 1933 2.147613e-04
55 1934 2.110177e-04
56 1935 2.144944e-04
57 1936 2.181519e-04
58 1937 2.169611e-04
59 1938 2.151304e-04
60 1939 2.179173e-04
61 1940 2.170693e-04
62 1941 2.144655e-04
63 1942 2.071141e-04
64 1943 2.076156e-04
65 1944 2.132730e-04
66 1945 2.163425e-04
67 1946 2.017719e-04
68 1947 1.889952e-04
69 1948 1.912796e-04
70 1949 1.907812e-04
71 1950 1.901058e-04
72 1951 1.873256e-04
73 1952 1.840656e-04
74 1953 1.808511e-04
75 1954 1.787660e-04
76 1955 1.764225e-04
77 1956 1.729511e-04
78 1957 1.694996e-04
79 1958 1.700200e-04
80 1959 1.663820e-04
81 1960 1.641475e-04
82 1961 1.605969e-04
83 1962 1.601109e-04
84 1963 1.590282e-04
85 1964 1.574027e-04
86 1965 1.629578e-04
87 1966 1.600521e-04
88 1967 1.566016e-04
89 1968 1.498546e-04
90 1969 1.407308e-04
91 1970 1.305796e-04
92 1971 1.256902e-04
93 1972 1.240797e-04
94 1973 1.214376e-04
95 1974 1.170620e-04
96 1975 1.119282e-04
97 1976 1.088652e-04
98 1977 1.041740e-04
99 1978 1.038398e-04
100 1979 9.941684e-05
101 1980 9.742368e-05
102 1981 9.729416e-05
103 1982 9.623298e-05
104 1983 9.768859e-05
105 1984 9.718731e-05

```
## 106 1985 9.424018e-05
## 107 1986 9.142422e-05
## 108 1987 8.806729e-05
## 109 1988 8.415047e-05
## 110 1989 7.908004e-05
## 111 1990 7.599062e-05
## 112 1991 7.469007e-05
## 113 1992 7.359058e-05
## 114 1993 7.190299e-05
## 115 1994 7.169631e-05
## 116 1995 7.137700e-05
## 117 1996 7.037749e-05
## 118 1997 6.878542e-05
## 119 1998 6.645501e-05
## 120 1999 6.489347e-05
## 121 2000 6.215487e-05
## 122 2001 6.105816e-05
## 123 2002 6.050068e-05
## 124 2003 5.934303e-05
## 125 2004 5.769723e-05
## 126 2005 5.681246e-05
## 127 2006 5.417926e-05
## 128 2007 5.278704e-05
## 129 2008 5.255107e-05
## 130 2009 5.307453e-05
## 131 2010 5.400867e-05
## 132 2011 5.433279e-05
## 133 2012 5.460906e-05
## 134 2013 5.547192e-05
## 135 2014 5.565828e-05
## 136 2015 5.593594e-05
## 137 2016 5.612298e-05
## 138 2017 5.689792e-05
```

OUTPUTS:

My outputs are propM which holds each mean for each year without including duplicates. The first item would represent 1880 and so on.

Each_year holds the corresponding year for each mean in propM.

years was a way to access a specific year so I could create a temporary matrix to calculate to mean.

Both of these values were then placed in a dataframe to organize them.

BODY and Sequence:


The for loop iterates from 1:138 (length of propM which is how many non duplicate years there are)

It begins by creating a temporary matrix which stores only the data from a particular year (by subsetting). During each iteration, tempMatrix would be reassigned a subset of babynames that only contained the data of a particular year.

I then used the mean() function to calculate the mean prop of the matrix and then added it to prop via indexing. I then also added the current year to Each_year. The loop would repeat this process once it has reached 138/length of prop.

5. Is there a relationship between **year** and its mean **prop**? Make the appropriate graph to show this relationship and add custom axis labels. Describe in words any relationship you see.

```
with(propM_and_year, plot(year, mean_prop))
```



A: mean_prop = average of proportions of each name in a given year (proportion of people of a particular sex with a particular name) year = each year that this study was conducted.

As the year increases, the prop of names overall decreases each year. A possible conclusion that can be drawn from this data is that people are having less children which is why names are becoming “less popular”. If there are less children being had then there would be less names being given, thus decreasing the prop overall each year.

6. Explain why a `for` loop is not the best way to complete the above task. Explore the `tapply` function and use it to produce the same variable you did in question 4. You could use `dplyr` functionality, as well, we just haven’t gotten there in class, yet.

```
tapply(babynames$prop, babynames$year, mean)
```

```
##      1880      1881      1882      1883      1884      1885
## 9.327296e-04 9.615283e-04 8.762491e-04 8.952682e-04 8.141704e-04 8.139021e-04
##      1886      1887      1888      1889      1890      1891
## 7.820985e-04 7.867618e-04 7.065885e-04 7.226397e-04 6.944554e-04 7.028995e-04
##      1892      1893      1894      1895      1896      1897
## 6.410328e-04 6.616759e-04 6.363961e-04 6.141118e-04 6.051479e-04 6.169039e-04
##      1898      1899      1900      1901      1902      1903
## 5.737293e-04 6.126111e-04 5.017828e-04 5.907095e-04 5.550258e-04 5.499067e-04
##      1904      1905      1906      1907      1908      1909
```



```
## 5.243434e-04 5.102156e-04 5.137318e-04 4.736429e-04 4.655718e-04 4.425246e-04
##      1910      1911      1912      1913      1914      1915
## 4.053799e-04 3.868689e-04 2.997358e-04 2.740336e-04 2.403288e-04 2.056436e-04
##      1916      1917      1918      1919      1920      1921
## 1.986359e-04 1.943714e-04 1.855249e-04 1.858838e-04 1.794639e-04 1.779018e-04
##      1922      1923      1924      1925      1926      1927
## 1.794109e-04 1.814584e-04 1.778298e-04 1.817062e-04 1.848431e-04 1.858832e-04
##      1928      1929      1930      1931      1932      1933
## 1.904784e-04 1.971559e-04 1.978530e-04 2.082921e-04 2.063539e-04 2.147613e-04
##      1934      1935      1936      1937      1938      1939
## 2.110177e-04 2.144944e-04 2.181519e-04 2.169611e-04 2.151304e-04 2.179173e-04
##      1940      1941      1942      1943      1944      1945
## 2.170693e-04 2.144655e-04 2.071141e-04 2.076156e-04 2.132730e-04 2.163425e-04
##      1946      1947      1948      1949      1950      1951
## 2.017719e-04 1.889952e-04 1.912796e-04 1.907812e-04 1.901058e-04 1.873256e-04
##      1952      1953      1954      1955      1956      1957
## 1.840656e-04 1.808511e-04 1.787660e-04 1.764225e-04 1.729511e-04 1.694996e-04
##      1958      1959      1960      1961      1962      1963
## 1.700200e-04 1.663820e-04 1.641475e-04 1.605969e-04 1.601109e-04 1.590282e-04
##      1964      1965      1966      1967      1968      1969
## 1.574027e-04 1.629578e-04 1.600521e-04 1.566016e-04 1.498546e-04 1.407308e-04
##      1970      1971      1972      1973      1974      1975
## 1.305796e-04 1.256902e-04 1.240797e-04 1.214376e-04 1.170620e-04 1.119282e-04
##      1976      1977      1978      1979      1980      1981
## 1.088652e-04 1.041740e-04 1.038398e-04 9.941684e-05 9.742368e-05 9.729416e-05
##      1982      1983      1984      1985      1986      1987
## 9.623298e-05 9.768859e-05 9.718731e-05 9.424018e-05 9.142422e-05 8.806729e-05
##      1988      1989      1990      1991      1992      1993
## 8.415047e-05 7.908004e-05 7.599062e-05 7.469007e-05 7.359058e-05 7.190299e-05
##      1994      1995      1996      1997      1998      1999
## 7.169631e-05 7.137700e-05 7.037749e-05 6.878542e-05 6.645501e-05 6.489347e-05
##      2000      2001      2002      2003      2004      2005
## 6.215487e-05 6.105816e-05 6.050068e-05 5.934303e-05 5.769723e-05 5.681246e-05
##      2006      2007      2008      2009      2010      2011
## 5.417926e-05 5.278704e-05 5.255107e-05 5.307453e-05 5.400867e-05 5.433279e-05
##      2012      2013      2014      2015      2016      2017
## 5.460906e-05 5.547192e-05 5.565828e-05 5.593594e-05 5.612298e-05 5.689792e-05
```

```
#system.time(tapply(babynames$prop, babynames$year, mean)) #about 3 second run time
```

A: For loops in general take longer in R since it is a statistical programming language. In this case, my loop ran faster than tapply function. I think for a much larger data set this will not be the case. Although it ran somewhat faster, it is more complicated than just using vectorize function. There are more variables, and more operations being performed (that are visible) which can make it harder to understand and make it slower when the ranges become larger (e.g 1- 1e⁹).

Bonus: Given my kids recent birthday, tell me your best dad joke. Double bonus if it has relevance to this class: